Priya Sharma

New York, NY | priya.sharma.ml@gmail.com | (212) 555-0144 | linkedin.com/in/priyasharma-ai | github.com/priyasharma-ml

Professional Summary

Machine Learning Engineer with 4+ years of experience specializing in Generative AI, Large Language Models (LLMs), RAG pipelines, and PyTorch. Proven ability to take research models into high-availability production environments using FastAPI, Docker, and AWS.

Technical Skills

Languages & Frameworks: Python, PyTorch, TensorFlow, Scikit-Learn, FastAPI, LangChain, LlamaIndex, HuggingFace
Core Competencies: Large Language Models (LLMs), RAG, NLP, Computer Vision, Transformers, Vector Search, Fine-Tuning
Cloud & Infrastructure: AWS, Docker, Kubernetes, Linux, Git, CI/CD, Redis, PostgreSQL, Pinecone, Milvus

Professional Experience

AI Systems Engineer | Apex AI Labs (2022 - Present)

- Built low-latency Retrieval-Augmented Generation (RAG) pipelines using LangChain, Qdrant vector database, and FastAPI serving over 500k queries/day.
- Fine-tuned open-source LLMs (Llama 3, Mistral) for domain-specific medical document summarization, achieving 92% factual accuracy.
- Deployed model inference endpoints on AWS SageMaker with Docker containerization and auto-scaling.

Machine Learning Engineer | DataCore Analytics (2020 - 2022)

- Developed NLP classification and sentiment models using PyTorch, HuggingFace Transformers, and Scikit-Learn.
- Automated training and evaluation workflows using Python, Pandas, and GitHub Actions CI/CD.
- Partnered with cross-functional teams to deploy real-time analytics features.

Education

Master of Science in Data Science | Columbia University (2018 - 2020)
Bachelor of Technology in Computer Science | Delhi Technological University (2014 - 2018)

Certifications

- AWS Certified Machine Learning - Specialty
- Deep Learning Specialization (DeepLearning.AI)